

Lecture 11: Empirical Bayes

*Instructor: Ryan Brill**Scribe: Jonathan Pipping*

11.1 Motivating Example: MLB Batting Averages

11.1.1 Single-Player Problem

Suppose player A 's batting average midway through the 2023 season is 0.300. Using no other information, we want to predict his end-of-season batting average. We set up a model to do this.

Let N and H represent the number of at-bats and hits for player A to this point in the season. We assume each at-bat $\{X_i\}_{i=1}^N$ is an independent Bernoulli trial with success probability p . Then it follows that

$$H \sim \text{Binomial}(N, p)$$

and since $H = \sum_{i=1}^N X_i$, we model player A 's mid-season batting average as

$$BA = \frac{H}{N} \sim \frac{1}{N} \text{Binomial}(N, p)$$

As discussed previously, the MLE of a binomial random variable is the observed proportion of successes, which we write as

$$\hat{p}_{MLE} = \frac{H}{N}$$

We can use this for prediction, but we saw that this may not be the most stable estimator in small samples. In the previous lecture, we used a prior to stabilize the estimator, but without any other information, we may not know which prior to use. Let's consider a broader approach to the problem that will allow us to inform our choice of prior.

11.1.2 Multi-Player Problem

Suppose now we know each player's batting average midway through the 2023 season. Using no information from any previous season (i.e. using only the 2023 mid-season averages), predict each player's end-of-season batting average.

We can easily extend our single-player model to the multi-player case. We set up notation for the multi-player case as follows:

- i : player index
- n : total number of players
- N_i : number of at-bats for player i
- H_i : number of hits for player i

- $BA_i = \frac{H_i}{N_i}$: batting average for player i

We can then model each player's batting average as before:

$$BA_i = \frac{H_i}{N_i} \sim \frac{1}{N_i} \text{Binomial}(N_i, p_i)$$

If we use the MLE for each player's batting average, we arrive at a similar result as before:

$$\hat{p}_i^{(MLE)} = \frac{H_i}{N_i}$$

Our hope is to improve this guess with a prior, but how do we do that? We don't have any other information, so we can't use a prior from a previous season. What can/should we shrink to?

11.1.3 Idea: Shrinkage by Pooling

Since we have the midseason batting average of **each** baseball player, perhaps we can pool information across these players and shrink to the overall mean batting average. *The insight here is that each player i is a baseball player, and we can use that shared group information to improve our prediction.* This means we would predict player i 's end-of-season batting average as some mixture of their own mid-season batting average and the overall mean batting average. Let's modify our model to reflect this.

Let $X_i = \frac{H_i}{N_i}$ be player i 's mid-season batting average. Then

$$X_i = \frac{H_i}{N_i} \sim \frac{1}{N_i} \text{Binomial}(N_i, p)$$

If we remove players i with a small number of at-bats (e.g. $N_i < 30$), we can use the fact that N_i are large to apply the Central Limit Theorem from Lecture 8. That is,

$$X_i \overset{approx}{\sim} \mathcal{N}\left(p_i, \frac{p_i(1-p_i)}{N_i}\right) \text{ by the CLT}$$

Note that the proposed variance $\sigma_i^2 = \frac{p_i(1-p_i)}{N_i}$ depends on the unknown parameter p_i , but it is much easier to work with **known** variance. We make the following *simplifying assumption* that

$$\sigma_i^2 = \frac{C}{N_i}$$

for some constant C . We'll use $C = 0.035$ for our analysis, but we generally treat this as a hyperparameter to be tuned using data from the previous season. An alternative approach is to use a **variance-stabilizing transformation** $h(X)$, which transforms the batting average so that it has a known variance. We will discuss this at a later time.

In either case, we have that $X_i \sim \mathcal{N}(\mu_i, \sigma_i^2)$ where $\mu_i = p_i$ and σ_i^2 is known. Now we can set up a parametric Bayesian model to predict each player's end-of-season batting average.

11.2 Parametric Bayesian Model

11.2.1 Model Setup

We implement the Bayesian approach, modeling the unknown parameter $\mu_i = p_i$ as a random variable with its own distribution. We then set up our model as follows:

$$\begin{cases} X_i \sim \mathcal{N}(\mu_i, \sigma_i^2) \\ \mu_i \sim \mathcal{N}(\mu, \tau^2) \end{cases}$$

Note that μ represents the global mean of all baseball players' batting averages, and τ^2 represents the variance around this mean. Note also that μ_i, p_i, μ , and τ are all unknown parameters, while $\sigma_i = \sqrt{\frac{C}{N_i}}$ is assumed as known (from before).

11.2.2 MLE vs Posterior Mean

Recall from before that the MLE ignores all information from other players, instead only using the information relevant to player i . Specifically, we know that the MLE of μ_i is the observed batting average X_i .

$$\hat{\mu}_i^{(MLE)} = X_i$$

The Bayesian analog of this quantity is the **posterior mean**, which we define below.

Definition 11.1 (Posterior Mean). *The posterior mean of a random variable X is defined as*

$$\hat{\theta}_{Bayes} = \mathbb{E}[\theta|X] = \int \theta p(\theta|X) d\theta$$

In our case, the posterior mean of μ_i from the data $\{X_i\}_{i=1}^n$ is

$$\hat{\mu}_i^{(Bayes)} = \mathbb{E}[\mu_i|X_i]$$

To go further than this, we need the posterior distribution $\mathbb{P}(\mu_i|X_i)$. We can derive this using Bayes' rule:

$$\begin{aligned} \mathbb{P}(\mu_i|X_i) &= \frac{\mathbb{P}(X_i|\mu_i)\mathbb{P}(\mu_i)}{\mathbb{P}(X_i)} \text{ by Bayes' Rule} \\ &\propto \mathbb{P}(X_i|\mu_i)\mathbb{P}(\mu_i) \text{ since } \mathbb{P}(X_i) \text{ is constant w.r.t. } \mu_i \end{aligned}$$

From our model setup, we re-express the likelihood and prior as

$$\begin{aligned} \mathbb{P}(\mu_i|X_i) &= \mathbb{P}(\mathcal{N}(\mu_i, \sigma_i^2) = X_i) \mathbb{P}(\mathcal{N}(\mu, \tau^2) = \mu_i) \\ &= \frac{1}{\sqrt{2\pi\sigma_i^2}} \exp\left(-\frac{(X_i - \mu_i)^2}{2\sigma_i^2}\right) \frac{1}{\sqrt{2\pi\tau^2}} \exp\left(-\frac{(\mu_i - \mu)^2}{2\tau^2}\right) \\ &\propto \exp\left[-\frac{1}{2}\left(\frac{X_i^2}{\sigma_i^2} - 2\frac{X_i\mu_i}{\sigma_i^2} + \frac{\mu_i^2}{\sigma_i^2} + \frac{\mu_i^2}{\tau^2} - 2\frac{\mu_i\mu}{\tau^2} + \frac{\mu^2}{\tau^2}\right)\right] \end{aligned}$$

Since the first and last terms are constant w.r.t. μ_i , we omit them and factor to get

$$\begin{aligned}\mathbb{P}(\mu_i|X_i) &\propto \exp \left\{ -\frac{1}{2} \left[\mu_i^2 \left(\frac{1}{\sigma_i^2} + \frac{1}{\tau^2} \right) - 2\mu_i \left(\frac{X_i}{\sigma_i^2} + \frac{\mu}{\tau^2} \right) \right] \right\} \\ &\propto \exp \left\{ -\frac{1}{2} \left(\frac{1}{\sigma_i^2} + \frac{1}{\tau^2} \right) \left[\mu_i^2 - 2\mu_i \frac{\left(\frac{X_i}{\sigma_i^2} + \frac{\mu}{\tau^2} \right)}{\left(\frac{1}{\sigma_i^2} + \frac{1}{\tau^2} \right)} \right] \right\} \\ &\propto \exp \left\{ -\frac{1}{2} \left(\frac{1}{\sigma_i^2} + \frac{1}{\tau^2} \right) \left[\mu_i - \frac{\left(\frac{X_i}{\sigma_i^2} + \frac{\mu}{\tau^2} \right)}{\left(\frac{1}{\sigma_i^2} + \frac{1}{\tau^2} \right)} \right]^2 \right\}\end{aligned}$$

We recognize this as a normal distribution, and so we can write

$$\mathbb{P}(\mu_i|X_i) \sim \mathcal{N} \left(\frac{\left(\frac{X_i}{\sigma_i^2} + \frac{\mu}{\tau^2} \right)}{\left(\frac{1}{\sigma_i^2} + \frac{1}{\tau^2} \right)}, \frac{1}{\left(\frac{1}{\sigma_i^2} + \frac{1}{\tau^2} \right)} \right)$$

Since we have a posterior distribution with a recognizable form, we can extract the posterior mean directly from the mean of the normal distribution. From Definition 11.1, we have that

$$\begin{aligned}\hat{\mu}_i^{(\text{Bayes})} &= \mathbb{E}[\mu_i|X_i] \\ &= \frac{\left(\frac{X_i}{\sigma_i^2} + \frac{\mu}{\tau^2} \right)}{\left(\frac{1}{\sigma_i^2} + \frac{1}{\tau^2} \right)} \\ &= \mu + \frac{\tau^2}{\tau^2 + \sigma_i^2} (X_i - \mu)\end{aligned}\tag{11.1}$$

And since we know that $X_i = \frac{H_i}{N_i}$ and $\sigma_i^2 = \frac{C}{N_i}$, we can write

$$\hat{\mu}_i^{(\text{Bayes})} = \frac{\frac{H_i}{C} + \frac{\mu}{\tau^2}}{\frac{N_i}{C} + \frac{1}{\tau^2}}\tag{11.2}$$

This looks a lot like the formulation we saw in the last lecture:

$$\frac{(W + W')}{(W + W') + (L + L')}$$

There's still a problem though: we don't know what μ or τ are! How do we estimate these parameters? Empirical Bayes offers a solution to this problem.

11.3 Empirical Bayes

Empirical Bayes is a way to estimate the parameters of a Bayesian model (in our case μ and τ) using the data itself. Specifically, this involves plugging in the MLEs of μ and τ into Equation 11.1, which we do to

define the **Parametric Empirical Bayes Estimator** for our model.

$$\begin{aligned}\hat{\mu}_i^{(EB)} &= \hat{\mu} + \frac{\hat{\tau}^2}{\hat{\tau}^2 + \sigma_i^2} (X_i - \hat{\mu}) \\ &= \frac{\frac{H_i}{C} + \frac{\hat{\mu}}{\hat{\tau}^2}}{\frac{N_i}{C} + \frac{1}{\hat{\tau}^2}}\end{aligned}\tag{11.3}$$

where $\hat{\mu}$ and $\hat{\tau}^2$ are the MLEs of μ and τ^2 . Before we derive these estimate, let's consider the implications of different values of τ^2 on the posterior mean.

11.3.1 Implications of Different τ^2 Values on Posterior Mean

If $\tau^2 = 0$, then $\mu_i \sim \mathcal{N}(\mu, \tau^2) \stackrel{d}{=} \mathcal{N}(\mu, 0) = \mu$.

If $\tau^2 = \infty$, then $\mu_i \sim \mathcal{N}(\mu, \tau^2) \stackrel{d}{=} \text{Uniform}(-\infty, \infty)$ and $\hat{\mu}_i = X_i = \frac{H_i}{N_i} = \hat{\mu}_i^{(MLE)}$.

Otherwise, we have that $\hat{\mu}_i = \hat{\mu} + \frac{\hat{\tau}^2}{\hat{\tau}^2 + \sigma_i^2} (X_i - \hat{\mu})$.

- Is closer to $\hat{\mu}$ if $\frac{\hat{\tau}^2}{\hat{\tau}^2 + \sigma_i^2}$ is small (meaning σ_i^2 is large, or N_i is small).
- Is closer to X_i if $\frac{\hat{\tau}^2}{\hat{\tau}^2 + \sigma_i^2}$ is large (meaning σ_i^2 is small, or N_i is large).

Now we will proceed by deriving $\hat{\mu}_{MLE}$ and $\hat{\tau}_{MLE}^2$.

11.3.2 Deriving $\hat{\mu}_{MLE}$ and $\hat{\tau}_{MLE}^2$

Recall that our model is

$$\begin{cases} X_i \sim \mathcal{N}(\mu_i, \sigma_i^2) \\ \mu_i \sim \mathcal{N}(\mu, \tau^2) \end{cases}$$

Since we want to find the MLE of μ and τ^2 , we need to maximize the likelihood function $\mathcal{L}(\mu, \tau^2 \mid X_i)$. Like we did in Lecture 10, we will instead maximize the log-likelihood function $\ell(\mu, \tau^2 \mid X_i)$.

$$\ell(\mu, \tau^2 \mid X_i) = \log \mathbb{P}(X_1, \dots, X_n \mid \mu, \tau^2)$$

Since the X_i are independent, we have that

$$\begin{aligned}\ell(\mu, \tau^2 \mid X_i) &= \log \prod_{i=1}^n \mathbb{P}(X_i \mid \mu, \tau^2) \\ &= \sum_{i=1}^n \log \mathbb{P}(X_i \mid \mu, \tau^2)\end{aligned}$$

From Bayes' Rules, we have that $\mathbb{P}(X_i | \mu, \tau^2) \sim \mathcal{N}(\mu, \tau^2 + \sigma_i^2)$. So we have that

$$\begin{aligned}\ell(\mu, \tau^2 | X_i) &= \sum_{i=1}^n \log \mathbb{P}(\mathcal{N}(\mu, \tau^2 + \sigma_i^2) = X_i) \\ &= \sum_{i=1}^n \log \left[\frac{1}{\sqrt{2\pi(\tau^2 + \sigma_i^2)}} \exp \left(-\frac{(X_i - \mu)^2}{2(\tau^2 + \sigma_i^2)} \right) \right] \\ &= \sum_{i=1}^n \left\{ \log \left(\frac{1}{\sqrt{2\pi(\tau^2 + \sigma_i^2)}} \right) - \frac{(X_i - \mu)^2}{2(\tau^2 + \sigma_i^2)} \right\} \\ &\propto -\frac{1}{2} \sum_{i=1}^n \log(\tau^2 + \sigma_i^2) - \frac{1}{2} \sum_{i=1}^n \frac{(X_i - \mu)^2}{\tau^2 + \sigma_i^2}\end{aligned}$$

Then we have that

$$MLE(\mu, \tau^2) = \arg \max_{\mu, \tau^2} \ell(\mu, \tau^2 | X_i)$$

which we can solve by taking the derivative with respect to μ and τ^2 and setting them equal to 0. First, for μ :

$$\begin{aligned}\frac{\partial}{\partial \mu} [\ell(\mu, \tau^2 | X_i)] &= \frac{1}{2} \sum_{i=1}^n \frac{2(X_i - \mu)}{\tau^2 + \sigma_i^2} = 0 \\ \implies \hat{\mu}_{MLE} &= \frac{\sum_{i=1}^n \frac{X_i}{\tau^2 + \sigma_i^2}}{\sum_{i=1}^n \frac{1}{\tau^2 + \sigma_i^2}}\end{aligned}$$

Now we take the derivative with respect to τ^2 :

$$\begin{aligned}\frac{\partial}{\partial \tau^2} [\ell(\mu, \tau^2 | X_i)] &= -\frac{1}{2} \sum_{i=1}^n \frac{1}{\tau^2 + \sigma_i^2} + \frac{1}{2} \sum_{i=1}^n \frac{(X_i - \mu)^2}{(\tau^2 + \sigma_i^2)^2} = 0 \\ \implies \sum_{i=1}^n \frac{(X_i - \mu)^2}{(\hat{\tau}_{MLE}^2 + \sigma_i^2)^2} &= \sum_{i=1}^n \frac{1}{\hat{\tau}_{MLE}^2 + \sigma_i^2}\end{aligned}$$

Now, we have a slight problem: our expression for $\hat{\mu}_{MLE}$ depends on τ^2 , and our expression for $\hat{\tau}_{MLE}^2$ depends on $\hat{\mu}$. We can solve this by iterating.

1. Make initial guesses $\mu_{(0)}$ and $\tau_{(0)}^2$.
2. While $|\mu_{(t)} - \mu_{(t-1)}| \leq \delta$ and $|\tau_{(t)}^2 - \tau_{(t-1)}^2| \leq \delta$, do the following in R:

$$\text{Set } \hat{\mu}_{(t)} = \frac{\sum_{i=1}^n \frac{X_i}{\tau_{(t-1)}^2 + \sigma_i^2}}{\sum_{i=1}^n \frac{1}{\tau_{(t-1)}^2 + \sigma_i^2}}$$

$$\text{Set } \hat{\tau}_{(t)}^2 \text{ to the solution of } \sum_{i=1}^n \frac{(X_i - \hat{\mu}_{(t)})^2}{(\hat{\tau}_{(t-1)}^2 + \sigma_i^2)^2} = \sum_{i=1}^n \frac{1}{\hat{\tau}_{(t-1)}^2 + \sigma_i^2} \text{ using the } \texttt{uniroot} \text{ function}$$

3. Iterate until convergence to $\hat{\mu}_{MLE}$ and $\hat{\tau}_{MLE}^2$.

11.3.3 Results

We plot the predicted batting averages for the 2023 season ($\hat{\mu}_i$) in Figure 11.1, with the MLE on the x-axis and the Empirical Bayes estimate on the y-axis. Note that the size of the points is proportional to the number of at-bats (N_i) for each player.




Figure 11.1: Comparison of MLE and EB estimates of 2023 end-of-season batting averages. The diagonal represents the identity line $EB = MLE$.

As we can see, players with smaller N_i have $\hat{\mu}_i^{(EB)}$ shrunk towards the overall mean, while players with larger N_i have $\hat{\mu}_i^{(EB)} \approx \hat{\mu}_i^{(MLE)}$, their mid-season batting averages. When we plot both estimates against the end-of-season batting averages in Figure 11.2, we don't see much of a difference between the two. However, considering their respective RMSE's in Figure 11.3, we see that the Empirical Bayes estimate achieves better prediction than the actual mid-season batting averages (MLE).




Figure 11.2: Comparison of MLE and EB estimates of end-of-season batting averages against the actual 2023 end-of-season batting averages. The diagonal represents the identity line $BA = \hat{BA}$.

rmse_MLE	rmse_EB
0.02629828	0.02383808

Figure 11.3: RMSE of the MLE and EB estimates of 2023 end-of-season batting averages.

11.3.4 Takeaways

- Shrinkage towards the overall mean helps prediction, especially when we have small sample sizes.
- Sharing information across players helps, and Empirical Bayes allows us to estimate these global parameters directly from the data.

11.3.5 If We Had Access to Previous Seasons

What if we had access to previous seasons' data? Consider our Parametric Empirical Bayes Estimator from Equation 11.3:

$$\hat{\mu}_i^{(EB)} = \hat{\mu} + \frac{\hat{\tau}^2}{\hat{\tau}^2 + \sigma_i^2} (X_i - \hat{\mu})$$

where $\hat{\mu}$ and $\hat{\tau}^2$ are the MLEs of μ and τ^2 . This is one example of a class of shrinkage estimators, which we will define generally now.

Definition 11.2 (General Shrinkage Estimator). *A general shrinkage estimator $\hat{\theta}_i$ takes the form*

$$\hat{\theta}_i = \hat{\theta} + \beta(X_i - \hat{\theta})$$

where $\hat{\theta}$ is the baseline (often the MLE or mean), $\beta \in [0, 1]$ is the shrinkage parameter, and X_i is the observed data.

The procedure to incorporate data from previous seasons is as follows:

1. Let $\{X_i^{(2022)}\}_{i=1}^n$ be the observed mid-season batting averages from 2022, let $\hat{\mu}^{(2022)}$ be the overall mean batting average from that season, and let $\hat{\mu}_i^{(2022)}$ be the known end-of-season batting averages from that year.
2. Estimate β using the following regression:

$$\begin{aligned} \hat{\mu}_i^{(2022)} &= \hat{\mu}^{(2022)} + \beta \left(X_i^{(2022)} - \hat{\mu}^{(2022)} \right) + \epsilon_i \\ \implies \hat{\mu}_i^{(2022)} - \hat{\mu}^{(2022)} &= \beta \left(X_i^{(2022)} - \hat{\mu}^{(2022)} \right) + \epsilon_i \\ \implies Y_i &= \beta X_i + \epsilon_i \end{aligned}$$

where $Y_i = \hat{\mu}_i^{(2022)} - \hat{\mu}^{(2022)}$ and $X_i = X_i^{(2022)} - \hat{\mu}^{(2022)}$. Call this estimate $\hat{\beta}$.

3. Predict the end-of-season batting averages for the current season as follows:

$$\hat{\mu}_i^{(2023)} = \hat{\mu}^{(2023)} + \hat{\beta} \left(X_i^{(2023)} - \hat{\mu}^{(2023)} \right)$$