

Lecture 8: Central Limit Theorem & Binomial Confidence Intervals

Instructor: Jonathan Pipping

Author: Ryan Brill

8.1 Central Limit Theorem

Theorem 8.1 (Central Limit Theorem). Suppose $\{X_i\}_{i=1}^n$ are any collection of **independent and identically-distributed (i.i.d.)** random variables with mean $\mu = \mathbb{E}[X_i] < \infty$ and variance $\sigma^2 = \text{Var}[X_i] < \infty$.

Then as $n \rightarrow \infty$, their sum $S_n = \sum_{i=1}^n X_i$ and sample mean $\bar{X}_n = \frac{S_n}{n}$ converge to the standard normal distribution. Specifically,

$$\frac{S_n - n\mu}{\sigma\sqrt{n}} \xrightarrow{d} \mathcal{N}(0, 1)$$

$$\frac{\frac{S_n}{n} - \mu}{\frac{\sigma}{\sqrt{n}}} \xrightarrow{d} \mathcal{N}(0, 1)$$

Definition 8.2 (Convergence in Distribution). A sequence of random variables $\{X_n\}_{n=1}^\infty$ converges in distribution to a random variable X if for all x in the support of X ,

$$\lim_{n \rightarrow \infty} \mathbb{P}(X_n \leq x) = \mathbb{P}(X \leq x)$$

In other words, as $n \rightarrow \infty$, the distribution of X_n converges to the distribution of X .

What does this mean? If $Z \sim \mathcal{N}(0, 1)$, the Central Limit Theorem assures that

$$\mathbb{P}\left(a \leq \frac{S_n - n\mu}{\sigma\sqrt{n}} \leq b\right) \rightarrow \mathbb{P}(a \leq Z \leq b) \text{ as } n \rightarrow \infty$$

Tons of quantities in sports are the sum or mean of i.i.d. random variables, so this normal approximation becomes extremely useful!

8.2 The Binomial Parameter

8.2.1 Setting Up The Model

Player A shoots n free throws, whose results are given by a sequence of random variables $\{X_i\}_{i=1}^n$, where

$$X_i = \begin{cases} 1 & \text{if the } i^{\text{th}} \text{ free throw is made} \\ 0 & \text{if the } i^{\text{th}} \text{ free throw is missed} \end{cases}$$

Then $S_n = \sum_{i=1}^n X_i$ is the total number of free throws made by player A . We model S_n as follows:

$$S_n \sim \text{Binomial}(n, p)$$

where n is the number of free throws attempted and p is the probability of making a free throw. A Binomial random variable is the sum of n independent Bernoulli(p) random variables.

8.2.2 Estimating the Binomial Proportion

We want to estimate player A 's probability of making a free throw, p from the data $\{X_i\}_{i=1}^n$. Our “best guess” of p is

$$\hat{p} = \frac{S_n}{n} = \frac{1}{n} \sum_{i=1}^n X_i$$

In fact, this is the **Maximum Likelihood Estimator (MLE)** of p . You will prove this in the Homework. But how confident should we be in this estimate? We answer this question by constructing an asymptotic confidence interval for p .

8.2.3 Confidence Interval for the Binomial Proportion

Recall that each X_i is a Bernoulli(p) random variable representing the result of the i^{th} free throw. We know the mean and variance of a Bernoulli random variable:

$$\begin{aligned}\mu &= \mathbb{E}[X_i] = p \\ \sigma^2 &= \text{Var}[X_i] = p(1-p)\end{aligned}$$

Since $\{X_i\}_{i=1}^n$ are i.i.d., we can use the CLT to approximate the distribution of $S_n = \sum_{i=1}^n X_i$. We don't know the true variance, but we can estimate it by $\hat{p}(1-\hat{p})$. Then by the CLT,

$$\frac{\hat{p} - p}{\sqrt{\frac{\hat{p}(1-\hat{p})}{n}}} \xrightarrow{d} \mathcal{N}(0, 1) \text{ as } n \rightarrow \infty$$

Then letting z_q be the q quantile (or $100 \cdot q^{th}$ percentile) of the standard normal distribution, we have

$$\mathbb{P} \left(-z_{1-\alpha/2} \leq \frac{\hat{p} - p}{\sqrt{\frac{\hat{p}(1-\hat{p})}{n}}} \leq z_{1-\alpha/2} \right) \approx 1 - \alpha$$

Taking $\alpha = 0.05$, we have that $z_{0.975} \approx 1.96$. And then

$$\mathbb{P} \left(-1.96 \leq \frac{\hat{p} - p}{\sqrt{\frac{\hat{p}(1-\hat{p})}{n}}} \leq 1.96 \right) \approx 0.95$$

Rearranging, we get a 95% confidence interval for p :

$$\mathbb{P} \left(\hat{p} - 1.96 \sqrt{\frac{\hat{p}(1-\hat{p})}{n}} \leq p \leq \hat{p} + 1.96 \sqrt{\frac{\hat{p}(1-\hat{p})}{n}} \right) \approx 0.95$$

This is known as the 95% **Wald** Confidence Interval for the Binomial parameter p . We give a more general formulation below:

Definition 8.3 (Wald Confidence Interval for p). *The $(1 - \alpha) \cdot 100\%$ Wald Confidence Interval for the Binomial parameter p is given by*

$$\hat{p} \pm z_{1-\alpha/2} \sqrt{\frac{\hat{p}(1-\hat{p})}{n}}$$

8.2.4 M&M's Example

Now suppose you buy an M&M's bag with 56 M&M's, 14 of which are blue. Supposing the color of each M&M is independently drawn from some distribution, what is a 95% confidence interval for p_{blue} the probability the company makes an M&M blue?

Here, we have $n = 56$ and $\hat{p}_{\text{blue}} = \frac{\text{number of blue M\&M's}}{\text{total number of M\&M's}} = \frac{14}{56} = 0.25$. Then the 95% Wald Confidence Interval for p_{blue} is given by

$$\begin{aligned}\hat{p}_{\text{blue}} \pm z_{0.975} \sqrt{\frac{\hat{p}_{\text{blue}}(1 - \hat{p}_{\text{blue}})}{n}} &= 0.25 \pm 1.96 \sqrt{\frac{0.25(1 - 0.25)}{56}} \\ &= 0.25 \pm 0.1134, \text{ or } [0.1366, 0.3634]\end{aligned}$$

This interval is symmetric about $\hat{p}_{\text{blue}} = 0.25$ and has length $2 \times 0.1134 = 0.2268$: a fairly wide interval! What happens if we were to observe the same proportion of blue M&M's, but with a larger sample size? Let $n = 400$ and $\hat{p}_{\text{blue}} = 0.25$. Then the 95% Wald Confidence Interval for p_{blue} is given by

$$\begin{aligned}\hat{p}_{\text{blue}} \pm z_{0.975} \sqrt{\frac{\hat{p}_{\text{blue}}(1 - \hat{p}_{\text{blue}})}{n}} &= 0.25 \pm 1.96 \sqrt{\frac{0.25(1 - 0.25)}{400}} \\ &= 0.25 \pm 0.0424, \text{ or } [0.2076, 0.2924]\end{aligned}$$

This interval is much narrower than the previous one, with a length of only $2 \times 0.0424 = 0.0848$. In general, the width of the confidence interval scales with $O(n^{-1/2})$. We will now move on to interpreting the confidence interval.

8.3 Interpreting the Confidence Interval

8.3.1 Frequentist Interpretation

In the previous section, we used the CLT to construct the Wald Confidence Interval in Definition 8.3. We know from this that

$$\mathbb{P}\left(\hat{p} - z_{1-\alpha/2} \sqrt{\frac{\hat{p}(1-\hat{p})}{n}} \leq p \leq \hat{p} + z_{1-\alpha/2} \sqrt{\frac{\hat{p}(1-\hat{p})}{n}}\right) \approx 1 - \alpha$$

but what does this probability actually mean? Under the model

$$\begin{aligned}S_n &= \sum_{i=1}^n X_i \sim \text{Binomial}(n, p) \\ X_i &\stackrel{\text{i.i.d.}}{\sim} \text{Bernoulli}(p)\end{aligned}$$

p is an **unknown, fixed constant**. This is known as a **frequentist interpretation** of the parameter. What are the implications of this setup? Since p is fixed, the probability that the confidence interval contains p is *actually* either 0 or 1. So what does the probability from our confidence interval actually mean?

8.3.2 The Confidence Interval as a Random Variable

Believe it or not, **the confidence interval itself is a random variable**. Why? It depends on \hat{p} , which depends on random variables $\{X_i\}_{i=1}^n$ through S_n . So the confidence interval is a random variable, and the

probability that the confidence interval contains p is the probability that this random interval contains the fixed parameter p .

If we repeated the experiment many times, in each replication p remains the same, but the data $\{X_i\}_{i=1}^n$ changes by randomness, and by extension \hat{p} and the CI's also change. However, at our specified α level, we expect the CI to contain p in $100(1 - \alpha)\%$ of the replications.

8.3.3 Coverage

Definition 8.4 (Coverage). *The coverage of a confidence interval is the probability that the confidence interval contains the true parameter. Letting θ be the true parameter,*

$$\text{Coverage} = \mathbb{P}(\theta \in CI)$$

The Wald Confidence Interval is based on 2 approximations:

1. That $\mathbb{P}\left(p - z_{1-\alpha/2}\sqrt{\frac{p(1-p)}{n}} \leq p \leq p + z_{1-\alpha/2}\sqrt{\frac{p(1-p)}{n}}\right) \approx 1 - \alpha$ by the CLT
2. That we can plug in \hat{p} for p in $\sigma = \sqrt{\frac{p(1-p)}{n}}$

Because of these approximations, the **actual** coverage of the $100(1 - \alpha)\%$ Wald Confidence Interval can be quite far below the **nominal** coverage of $100(1 - \alpha)\%$ as shown by simulations and computations (Brown, Cai, Dasgupta, 2001). When does this happen?

If n is several hundred or thousand and/or p is close to 0.5, the Wald interval is generally **tolerably** accurate. However, if n is smaller or p is close to 0 or 1, the Wald interval can be quite inaccurate. To correct for this, Agresti and Coull (1998) recommend introducing **2 artificial successes and failures** into the data before computing \hat{p} , which is known as the **Agresti-Coull Interval**. We define this interval below.

Definition 8.5 (Agresti-Coull Interval). *The $(1 - \alpha) \cdot 100\%$ Agresti-Coull Interval for the Binomial parameter p is given by*

$$\hat{p} \pm z_{1-\alpha/2} \sqrt{\frac{\hat{p}'(1 - \hat{p}')}{n + 4}}, \text{ where } \hat{p}' = \frac{S_n + 2}{n + 4}$$

In many cases, the Agresti-Coull interval achieves much better coverage than the Wald interval. This is pictured in Figure 8.1, where the Agresti-Coull interval outperforms the Wald interval whenever p is even moderately far from 0.5.




Figure 8.1: Coverage of 95% Wald vs. Agresti-Coull Confidence Intervals

References

- [AC] Agresti, A., & Coull, B.A., *Approximate is Better than "Exact" for Interval Estimation of Binomial Proportions*, The American Statistician, 1998.
- [BCD] Brown, T.C., Cai, T.T., & Dasgupta, A., *Interval Estimation for a Binomial Proportion*, Statistical Science, 2001.